

When AI becomes the new gateway to your business

Written by Nikolai Waltham · Waltham Consulting
October 2026

Over the past year, I have been working increasingly closely with AI, particularly its impact on business visibility, customer journeys and the way companies operate.

AI visibility is already a field of work, with its own tools, specialists and terminology. The question that interests me is what a business can actually do with it.

Knowing that AI can recommend a supplier is one thing. Understanding why a customer asks a particular question, what evidence they need and how the business should respond takes a different kind of work.

I think this is where the architect's role becomes important: connecting technical capability to a customer need and a working business process.

Search remains part of that process. People still compare websites and decide whom to contact. But AI can now take on some of that comparison before the business ever hears from the customer.

Increasingly, customers will turn to their AI to find the right business:

"Who can help me with this?"

"Which three suppliers should I speak to?"

"What is the difference between them?"

"Who is the best fit for what I need?"

And if that AI already understands the person, their business, their budget and their circumstances, the question becomes even more specific.

That changes something fundamental.

Businesses are no longer competing only for a place at the top of the search results.

They are competing to be part of the answer at all.

If AI creates the shortlist

We saw a practical example of this with a bank we work with.

In our AI visibility checks, it did not appear when we asked a question along the lines of "I need a new bank for my business." It did appear when we narrowed the question to the particular category of bank it belonged to.

That distinction mattered. A business owner may describe what they need without knowing which category of bank to ask for.

Changes were made, and the bank subsequently appeared in our AI checks where it had previously been absent. That was an observed improvement; it did not establish which individual change caused it or guarantee inclusion in future answers.

The useful lesson was that being discoverable under your own industry terminology does not necessarily mean being found through your customer's language.

The customer should not have to know your category before they can discover your business.

If no one knows where you live, no one comes to visit.

The same principle applies online.

If AI does not know your business exists, or does not understand what you do well, you may never be included when a customer asks whom they should choose.

And if AI creates the customer's shortlist and you are not on it, you are out of the running before you even know the competition has begun.

The problem is not necessarily the website.

It may be that AI has too little evidence to understand the company's expertise, or to trust its claims.

From search rankings to a credible recommendation

SEO and AI visibility share important foundations: accessible websites, useful content and credible sources. AI-generated answers add a further challenge: how the business is represented when information from different sources is brought together.

For a business, the useful questions are:

Who is this company?

What is it an expert in?

Who says so?

Does the company's own description match what other sources say?

Do its people have genuine professional authority?

Is the company mentioned in relevant contexts?

Are there articles, customers, media coverage or other sources that substantiate its expertise?

Increasingly, this is about a company's digital reputation.

Perhaps we need to look back a few hundred years

The more I think about it, the more this development seems to take us back to something very old.

Imagine you needed to find a craftsperson 200 years ago. Who would you choose?

The person standing in the town square, shouting that they were the leading, most innovative and unquestionably best craftsperson in town?

Or someone recommended by three people you trusted? Someone whose work you had seen? Someone who could explain, calmly and precisely, what they would do, what it would cost and why they recommended that particular solution?

For centuries, we have judged credibility in much the same way: references, reputation, clarity, proven experience, other people's opinions and a sense that someone actually knows what they are talking about.

Yet so many company websites are still full of phrases such as:

"Market-leading." "Experts." "Unique." "Innovative." "Best in class." "We drive growth." "We are passionate."

When everyone says the same thing, those words begin to mean very little.

Perhaps we are moving in the opposite direction: less self-praise and more evidence. Less advertising language and greater precision.

Less "we are brilliant" and more:

"Here is what we do."

"This is how we do it."

"Here are some examples."

"And here are people who can vouch for us."

What interests me is that this probably works for both people and AI.

Digital reputation becomes more important

Any company can describe itself on its own website as a market leader, an innovator or one of the best in its industry.

The interesting part is when other sources tell the same story.

When employees share their professional knowledge. When customers talk about the company. When its expertise is quoted. When LinkedIn, industry publications and independent sources consistently associate it with the same area of expertise.

This is how people build trust.

Those independent sources can also give AI systems evidence to draw on when describing or recommending a business.

This is where market research matters again

In my view, this has another consequence that is easy to overlook.

If a website needs to do more than satisfy a search engine, and must persuade both people and AI that the business is relevant and credible, understanding the customer becomes even more important.

What are customers uncertain about?

What questions do they ask before buying?

What makes them nervous?

Which alternatives are they considering?

What do they need to know before they can trust a supplier?

What words do they use to describe their problem?

And what ultimately prompts them to pick up the phone, send an enquiry or click "buy"?

This is classic market research.

The bank example brings this into focus. A check based only on the bank's own category could have suggested that it was visible. Asking from the business owner's perspective exposed a gap.

That is why the practical work should begin with actual enquiries and customer conversations. What do people ask for? Which words do they use? With appropriate access, AI can help organise those questions, while someone who understands the business checks the interpretation and decides what needs to change.

You can have the most beautiful website in the world and the most advanced AI tools. But if you do not understand what is going on in your customer's mind, your communication will still fall short.

Technology changes quickly. The fundamentals of human decision-making change much more slowly.

That is why I believe customer understanding becomes more valuable in an AI world, not less.

The new risk is becoming invisible without knowing it

This is where I think many businesses underestimate what is happening.

A company can be highly capable, well established and competitive, yet never enter the customer's consideration at all.

You do not necessarily lose that customer to a better business. You lose them to the business AI actually knows.

Perhaps the most dangerous part is that you may never find out.

There is no rejection. No enquiry that fails to convert. No customer saying, "We chose someone else."

The customer simply never reaches you.

AI Visibility is already here

Businesses can already buy tools and services to track how they appear in AI answers. The questions are immediate:

"What does AI know about us?"

"What is AI telling our potential customers about us?"

But a visibility score only becomes useful when it leads to a sound decision. Which customer needs are we failing to address? What evidence is missing? Who will do the work, and how will we know whether it helped?

But why should people do all the work

If AI can analyse a company's digital visibility, why should people have to carry out the entire analysis themselves?

Parts of this work can be automated, provided the questions and the interpretation are grounded in the business.

A company enters its domain. The system identifies its market, products and relevant competitors. It tests relevant questions across different AI systems and examines the company's website, LinkedIn profiles, mentions and other available sources.

It can then answer:

How often is the company mentioned?

Which competitors are mentioned more often?

For which questions is the company being overlooked?

Which signals of credibility are missing?

What content should the company create?

And what specific changes should it make?

The measurement can be repeated month after month. Because AI answers vary, one response should never be treated as a definitive verdict. A useful service needs repeatable tests, the underlying answers and sources, and a clear account of uncertainty.

That creates scope for a recurring product. It also makes the design of that product important: a dashboard should help someone decide what to do next.

And this is where it gets really interesting

AI Visibility is probably only one of the first layers.

Businesses are already beginning to use specialised AI systems for different tasks: marketing, CRM, customer service, reporting, analysis, and their websites and digital visibility.

Another might handle accounting, with access to the company's financial system and the ability to work directly with its data and processes.

The next logical step goes beyond adding more isolated AI tools.

It is an overarching layer that can coordinate them: an AI administrator with a view across the business.

Consider a hypothetical situation: a company receives as many enquiries as before, but wins fewer orders. A marketing response might be to generate more leads. Yet the real problem could be that quotations take too long to reach customers.

With suitable access, a coordinating AI could ask a CRM agent to examine enquiry-to-quotation times and an operations agent to check capacity. It could bring the findings together for the sales manager to review. If delays are confirmed, an agreed workflow could flag overdue quotations, gather missing information and prepare follow-ups for approval.

The outcome to measure would be faster responses and a change in conversion, alongside the quality of the quotations. More AI activity would not, by itself, count as success.

The individual AI agents become the specialists. The overarching orchestration layer becomes the coordinator.

In my view, this is a much bigger development than simply getting a better chatbot.

AI is becoming a new operational management layer within the business.

Making that work requires decisions about data, responsibilities, permissions and when a person needs to intervene. It also requires someone to understand how a change in one part of the business affects another.

The architect has to understand the business

Technical expertise is essential. The architect's job is to give it direction: define the problem, connect the right capabilities and make the resulting workflow usable by the people responsible for it.

That role starts with questions such as: Where does work stall? What does the customer need? Which information can we trust? What may the system do on its own? Who takes over when it is wrong?

I believe the ability to answer those questions and turn them into working processes will increasingly distinguish businesses that create value with AI.

The value becomes visible when a customer gets a better answer, a quotation arrives sooner or a problem is resolved with less effort.

It will probably begin with small problems

No one needs to start with an AI that runs the entire business.

The useful starting point is a problem that is straightforward, measurable and relatively low risk.

AI Visibility could be one of those problems.

A clear need. A solution that can be automated. A service that almost any business with a website could potentially need.

And, at the same time, a practical way to learn how AI can become part of day-to-day operations.

Once businesses get used to AI doing the work as well as offering advice, expectations change quickly.

The question then shifts from:

“What can AI help us with?”

to:

“Which processes should people still be carrying out themselves?”

That is a much bigger discussion.

AI Visibility is unlikely to be the last product to emerge from this development. It could, however, be one of the first practical entry points into a business where AI gradually moves from being a supporting tool to becoming part of the operation itself.

And if the customer's AI also becomes the new gatekeeper to the market, one very simple question follows:

Does it even know your business exists?